

The product, *refocused*

One clear promise for 50 waitlist users: metered AI trading sessions that find a trade, and an always-on agent that manages it after entry. Free to run, honest about its edge, shippable now.

Doc	Date	Status	Supersedes
02 of 06 – Product plan	2026-08-24	Adopted direction	Tier table in docs/HLD.md §5

What CryptAI is — and the three promises it can keep

An AI trading co-pilot for crypto perpetual futures. The user brings an exchange account (or a paper desk); CryptAI brings a validated market-condition engine, a metered session where AI agents propose trades, and an unmetered monitor that manages every open position.

● PROMISE 1 · TIMING

"Is now a good time to trade?"

Answered by the Rust signal engine's **tradability score** — empirically validated as a short-horizon filter (IC +0.08–0.10), not vibes. Surfaced as a **hot-window banner**: spend your session slots when the market is actually tradable.

● PROMISE 2 · PERSONAL

"A trade team working your book."

Setups are synthesized **per user** — same chart, different trade — sized against your capital, risk appetite, monthly goal, and open positions. Every trade arrives as a proposal with entry, stop, targets, size, and a plain-language rationale. **You approve; it executes.**

● PROMISE 3 · VIGILANCE

"It doesn't fall asleep after entry."

On fill, a per-position **monitor agent** takes over — unmetered, outliving the session and every restart. It re-checks the thesis as conditions drift and decides: hold, tighten stop, take partial, or exit. **Finding trades is metered; protecting them never is.**

THE ONE-LINE PITCH

"Two 30-minute sessions a day with a professional AI trade desk — and an agent that watches every position you take, around the clock, for free."

Free tier: 2 × 30-minute sessions, one trade in focus

1

Hot window

The global engine scores market tradability continuously. Dashboard shows “**market is hot / quiet**” so users spend slots at the best time — no session needed to see it.



2

Start a session

One of **2 daily slots**, 30 min each, server-clocked. Agents scan and synthesize setups personalized to the user's capital, risk, and goal.



3

Proposal

A setup card arrives: entry / stop / targets / size / leverage / R:R + rationale, with a shelf life. **The clock pauses** while the user decides.



4

Trade focus

On approval the trade executes and the session goes **synchronous**: no new setups. The user's whole window pivots to the one running trade.



5

Managed exit

The monitor manages the position to profit-book or stop — even after the session ends. Portfolio and history update; the next session unlocks when the trade closes.

WHY SESSIONS, NOT ALWAYS-ON

- **Scarcity is the product:** session slots make the free tier generous but bounded, and the upgrade path obvious.
- **Cost follows sessions:** LLM spend only happens inside a metered window — the \$0 infra budget holds structurally.
- **Focus beats noise:** a 30-min window with a hot-market banner beats a firehose of all-day signals.

WHY SYNCHRONOUS ON FREE

- One open trade at a time = **capital protection by design** for beginners.
- Monitoring one trade is **cheap and unmetered** — the "never sleeps" promise costs us almost nothing.
- Async multi-trade becomes a **real, felt upgrade** at Lite/Pro instead of an arbitrary paywall.

SESSION RULES (FREE)

- 2 slots/day · 30 min each · server-authoritative clock · resets at 00:00 IST.
- Clock **pauses** while a proposal awaits a decision.
- While a trade is open, new sessions are **held** — slots are preserved, monitor keeps working, banner explains why.
- Unused slots don't roll over.

Four tiers — only Free ships now, the rest is roadmap

No purchasable pricing until a real payment rail exists (Razorpay/Stripe — roadmap M4). The landing page presents Free as a beta and the paid tiers as "coming soon". Env-var tier grants are replaced by a subscriptions table before any money moves.

CAPABILITY	FREE — SHIPS NOW	LITE	PRO	PRO MAX
Sessions / day	2 × 30 min	4 × 45 min	8 × 60 min	Unlimited windows
Trade mode	Synchronous — 1 open trade, session pivots to it	Async-lite — 2 concurrent trades	Async — 5 concurrent trades	Async — 10 concurrent trades
Channels	Intraday only	+ Swing	+ Scalp [†] , advanced setups	All + experimental strategies
Execution	Paper desk; exchange connect in testnet beta	+ Live (Delta India), consent-gated	+ BingX · auto-mode within consent caps	+ Multi-exchange, API access
Monitoring	Unmetered on every tier — position protection is never paywalled			
Assistant	—	Ask about your positions	Full market + book Q&A	Full + strategy tuning
Price	\$0	TBD · after payments ship	TBD	TBD

[†] Scalp channel carries an India honesty gate: 1% TDS per transfer materially eats scalping edge for Indian users — the UI must say so before a user picks it (see quant assessment, doc 01).

The honest edge — claims the data supports

● WE CLAIM — AND CAN DEFEND

- "Now is / isn't a good time to trade." Tradability score validated out-of-sample as a short-horizon filter (IC +0.083–0.096 at 4 bars; survives volatility normalization; re-validated on six held-out symbols).
- "This setup fits your book." Sizing and gating are deterministic against the user's real capital, risk %, and open positions.
- "Your position is being watched." Monitor re-evaluates on regime flip, pulse veto, and invalidation breach — mechanical, auditable triggers.

● WE NEVER CLAIM

- **Direction from regime.** Killed out-of-sample (walk-forward IC –0.008). No consumer may map regime → long/short. Contract-tested.
- "3 PM will be the best hour." The score is a **now-cast**, not a forecast. The hot-window banner says "market is hot now" and "typically active hours" — never a predicted golden hour, until logged calibration data earns it.
- **Guaranteed returns.** Monthly range is the user's **goal input** for risk shaping — never a promise in UI copy.

● LAUNCH POSTURE: PAPER-FIRST

Waitlist launch runs on paper desks (real fills simulation) + Delta India testnet connect. Real-money mainnet stays double-gated off until the compliance floor exists (ToS, risk disclosure, automated-execution consent, API-key scope verification, KYC). Rationale:

- 50 strangers × leveraged perps × an engine with a **brand-new wiring layer** = uninsurable risk on day one.
- Paper mode exercises 100% of the product loop — sessions, proposals, focus mode, monitoring, history.
- India: FIU-IND / VDA / advisory-licensing questions are open; paper keeps them dormant while we learn.

FALSIFIABILITY FROM DAY ONE

- Every published pulse is appended to the `pulse_snapshots` calibration log — the score stays testable against forward returns.
- Every proposal → outcome pair is recorded, so "did our setups make money" is a query, not a belief.
- Hot-window hours are learned from our own logged data before we ever present them as advice.

Waitlist launch: must / should / won't

● MUST — LAUNCH BLOCKS ON THESE

- Google sign-in + 4-step onboarding (mode → capital → risk → monthly goal)
- Hot-window banner + per-symbol pulse on dashboard
- 2x30-min session flow: start → setups → proposal card → approve → trade focus
- Paper execution + portfolio/history that survives restarts
- Per-position monitor + visible monitor window ("what my agent is doing")
- Quota integrity: server-clamped sessions, per-session LLM budget fed & enforced
- Security floor: admin fail-closed, no fail-open anon mode, destructive scripts out of the image

● SHOULD — FIRST WEEKS AFTER

- Delta India **testnet** live-connect beta (real order path, fake money)
- Learned "typically hot hours" from calibration log
- Session summary recap ("your 30 minutes: 3 setups, 1 trade, +\$142")
- Email/push nudge when the market turns hot and slots remain
- Waitlist referral loop ("skip the queue")

● WON'T — DELIBERATELY NOT NOW

- Payments, paid tiers, checkout — Free beta only
- Real-money mainnet execution — gated until compliance floor
- Scalp channel — TDS honesty + validation gap
- AI chat assistant — Lite/Pro roadmap
- CoinDCX — no testnet exists; unvalidatable without real money
- Mobile app — responsive web only

How we'll know it worked — and what only the founder can decide

SUCCESS METRICS — FIRST 30 DAYS		
METRIC	TARGET	WHY IT MATTERS
Waitlist activation	≥ 60%	50 invited → 30 complete onboarding + first session
Session utilization	≥ 1.0/day	Avg slots used per active user — the core habit
Proposal approval rate	25–60%	Below = setups aren't credible; above = users rubber-stamp
D7 retention	≥ 40%	Sessions are a daily mechanic — this is the product verdict
Monitor saves	count > 0	Exits triggered by thesis-break before SL — the differentiator, demo-able
LLM cost / session	\$0.00	Free-model routing holds; budget breaches alert, never bill

FOUNDER DECISIONS NEEDED (BLOCKING NONE OF THE BUILD)	
• Session reset timezone:	00:00 IST assumed — confirm.
• Paper capital:	seed at user-declared capital (realism) vs fixed \$10k (comparability). Plan assumes user-declared, capped \$1k–\$100k .
• Waitlist rollout:	all 50 at once vs 10/day cohorts. Plan assumes cohorts of 10 — engine + quota behavior observed under load before full open.
• Tier pricing:	deferred until payments ship — no number needed now.
• Brand voice on hot windows:	"market is hot" vs "conditions favorable" — copy choice, compliance-relevant.